

Chapter 10

Multiprocessor and Real-Time Scheduling

Agenda

■ Multiprocessor Scheduling

■ Real-Time Scheduling

- Linux Scheduling
- Linux Virtual Machine Process Scheduling

Agenda

■ Multiprocessor Scheduling

- Granularity
- Design Issues
- Process Scheduling
- Thread Scheduling

■ Real-Time Scheduling

- Background
- Characteristics of Real-Time Operating Systems
- Real-Time Scheduling
- Deadline Scheduling
- Rate Monotonic Scheduling
- Priority Inversion

Classifications of Multiprocessor Systems

Loosely coupled or distributed multiprocessor, or cluster

- consists of a collection of relatively autonomous systems, each processor having its own main memory and I/O channels

Functionally specialized processors

- there is a master, general-purpose processor; specialized processors are controlled by the master processor and provide services to it

Tightly coupled multiprocessor

- consists of a set of processors that share a common main memory and are under the integrated control of an operating system

Synchronization Granularity and Processes

Grain Size	Description	Synchronization Interval (Instructions)
Fine	Parallelism inherent in a single instruction stream.	<20
Medium	Parallel processing or multitasking within a single application	20-200
Coarse	Multiprocessing of concurrent processes in a multiprogramming environment	200-2000
Very Coarse	Distributed processing across network nodes to form a single computing environment	2000-1M
Independent	Multiple unrelated processes	not applicable

Independent Parallelism

- No explicit synchronization among processes
 - each represents a separate, independent application or job
- Typical use is in a time-sharing system

each user is performing a particular application

multiprocessor provides the same service as a multiprogrammed uniprocessor

because more than one processor is available, average response time to the users will be less



Coarse and Very Coarse-Grained Parallelism

- Synchronization among processes, but at a very gross level
- Good for concurrent processes running on a multiprogrammed uniprocessor
 - can be supported on a multiprocessor with little or no change to user software



Medium-Grained Parallelism

- Single application can be effectively implemented as a collection of threads within a single process
 - programmer must explicitly specify the potential parallelism of an application
 - there needs to be a high degree of coordination and interaction among the threads of an application, leading to a medium-grain level of synchronization
- Because the various threads of an application interact so frequently, scheduling decisions concerning one thread may affect the performance of the entire application



Fine-Grained Parallelism

- Represents a much more complex use of parallelism than is found in the use of threads
- Is a specialized and fragmented area with many different approaches



Design Issues

Scheduling on a multiprocessor involves three interrelated issues:

- The approach taken will depend on the degree of granularity of applications and the number of processors available

actual
dispatching
of a process

use of
multiprogramming on
individual processors

assignment of
processes to
processors

Assignment of Processes to Processors

Assuming all processors are equal, it is simplest to treat processors as a pooled resource and assign processes to processors on demand

static or dynamic needs to be determined

If a process is permanently assigned to one processor from activation until its completion, then a dedicated short-term queue is maintained for each processor


advantage is that there may be less overhead in the scheduling function

allows group or gang scheduling

- A disadvantage of static assignment is that one processor can be idle, with an empty queue, while another processor has a backlog
 - to prevent this situation, a common queue can be used
 - another option is dynamic load balancing

Assignment of Processes to Processors

- Both dynamic and static methods require some way of assigning a process to a processor
- Approaches:
 - Master/Slave
 - Peer



Master/Slave Architecture

- Key kernel functions always run on a particular processor
- Master is responsible for scheduling
- Slave sends service request to the master
- Is simple and requires little enhancement to a uniprocessor multiprogramming operating system
- Conflict resolution is simplified because one processor has control of all memory and I/O resources

Disadvantages:

- failure of master brings down whole system
- master can become a performance bottleneck

Peer Architecture


- Kernel can execute on any processor
- Each processor does self-scheduling from the pool of available processes

Complicates the operating system

- operating system must ensure that two processors do not choose the same process and that the processes are not somehow lost from the queue


Process Scheduling

- Usually processes are not dedicated to processors
- A single queue is used for all processors
 - if some sort of priority scheme is used, there are multiple queues based on priority
- System is viewed as being a multi-server queuing architecture



Process Scheduling

- With static assignment: should individual processors be multiprogrammed or should each be dedicated to a single process?
- Often it is best to have one process per processor; particularly in the case of multithreaded programs where it is advantageous to have all threads of a single process executing at the same time.



Thread Scheduling

- Thread execution is separated from the rest of the definition of a process
- An application can be a set of threads that cooperate and execute concurrently in the same address space
- On a uniprocessor, threads can be used as a program structuring aid and to overlap I/O with processing
- In a multiprocessor system kernel-level threads can be used to exploit true parallelism in an application
- Dramatic gains in performance are possible in multi-processor systems
- Small differences in thread management and scheduling can have an impact on applications that require significant interaction among threads

Approaches to Thread Scheduling

processes are not assigned to a particular processor

Load Sharing

a set of related thread scheduled to run on a set of processors at the same time, on a one-to-one basis

Gang Scheduling

Four approaches for multiprocessor thread scheduling and processor assignment are:

provides implicit scheduling defined by the assignment of threads to processors

Dedicated Processor Assignment

the number of threads in a process can be altered during the course of execution

Dynamic Scheduling


Load Sharing

- Simplest approach and carries over most directly from a uniprocessor environment

Advantages:

- load is distributed evenly across the processors
- no centralized scheduler required
- the global queue can be organized and accessed using any of the schemes discussed in Operating systems-1

- Versions of load sharing:
 - first-come-first-served
 - smallest number of threads first
 - preemptive smallest number of threads first



Disadvantages of Load Sharing

- Central queue occupies a region of memory that must be accessed in a manner that enforces mutual exclusion
 - can lead to bottlenecks
- Preemptive threads are unlikely to resume execution on the same processor
 - caching can become less efficient
- If all threads are treated as a common pool of threads, it is unlikely that all of the threads of a program will gain access to processors at the same time
 - the process switches involved may seriously compromise performance



Gang Scheduling

- Simultaneous scheduling of the threads that make up a single process

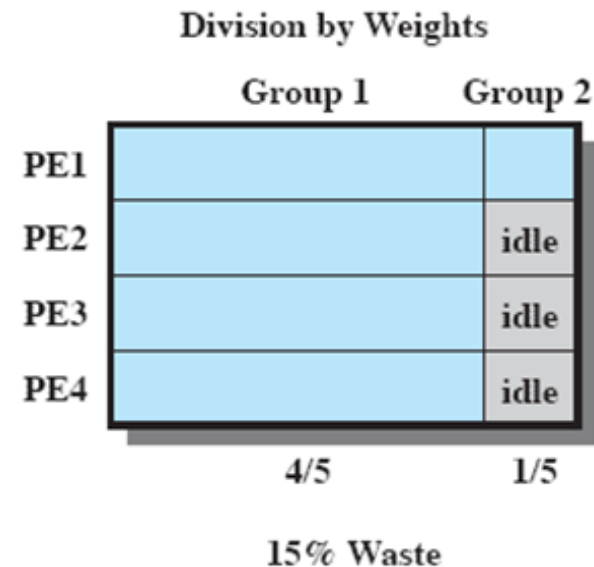

Benefits:

- synchronization blocking may be reduced, less process switching may be necessary, and performance will increase
- scheduling overhead may be reduced

- Useful for medium-grained to fine-grained parallel applications whose performance severely degrades when any part of the application is not running while other parts are ready to run
- Also beneficial for any parallel application

Figure 10.2

Example of Scheduling Groups With Four and One Threads



Dedicated Processor Assignment

- When an application is scheduled, each of its threads is assigned to a processor that remains dedicated to that thread until the application runs to completion
- If a thread of an application is blocked waiting for I/O or for synchronization with another thread, then that thread's processor remains idle
 - there is no multiprogramming of processors
- Defense of this strategy:
 - in a highly parallel system, with tens or hundreds of processors, processor utilization is no longer so important as a metric for effectiveness or performance
 - the total avoidance of process switching during the lifetime of a program should result in a substantial speedup of that program

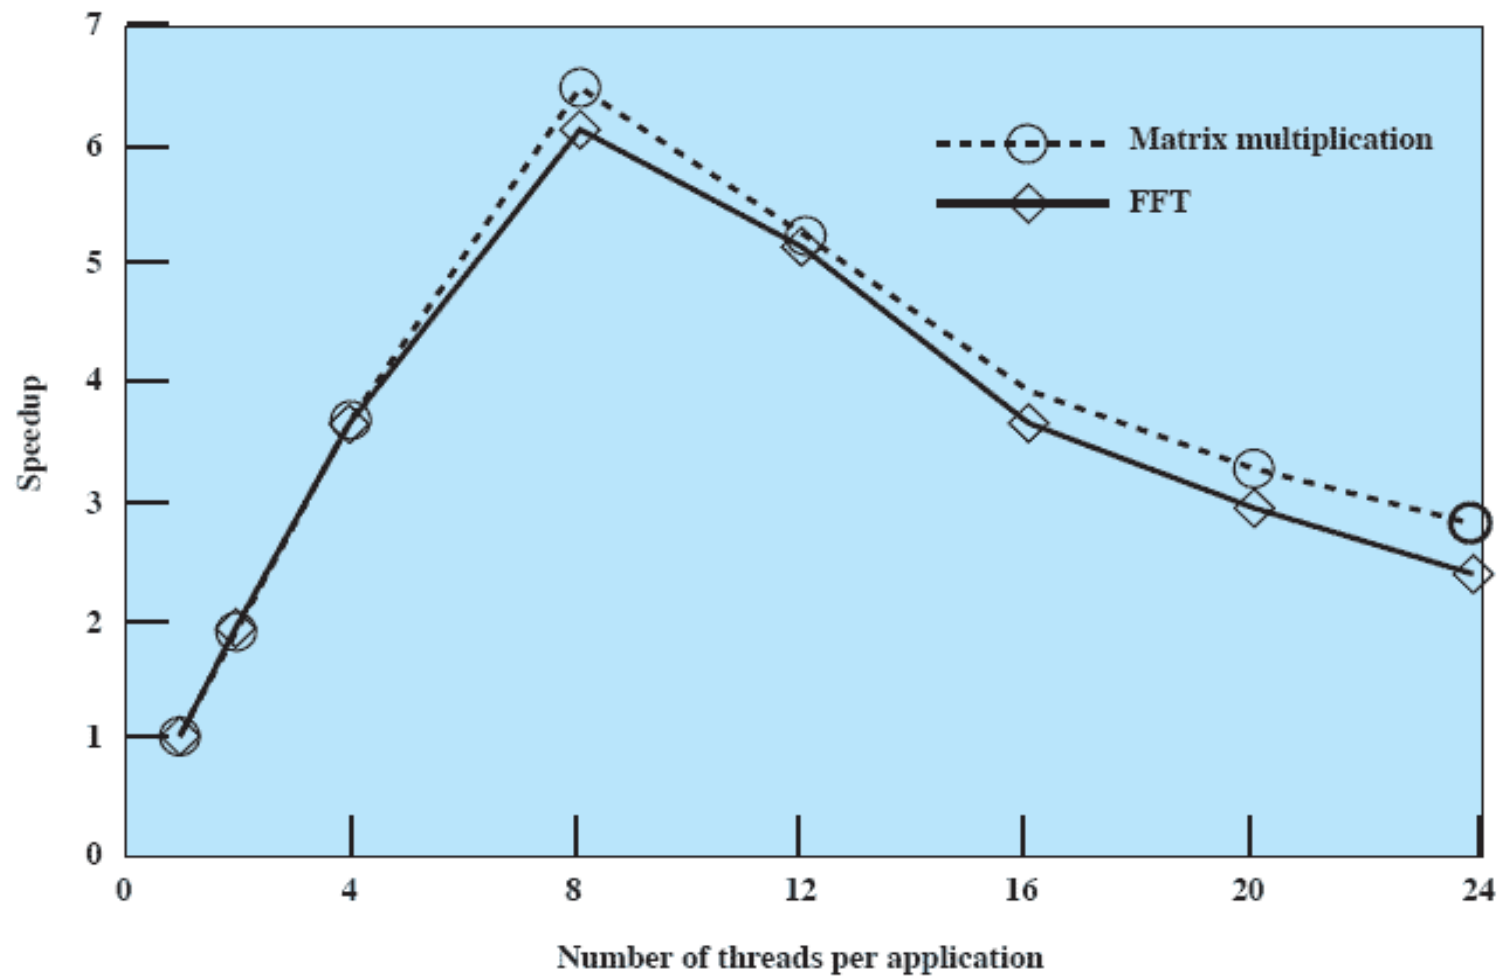


Figure 10.3

Application Speedup as a Function of Number of Threads

Dynamic Scheduling

- For some applications it is possible to provide language and system tools that permit the number of threads in the process to be altered dynamically
 - this would allow the operating system to adjust the load to improve utilization
- Both the operating system and the application are involved in making scheduling decisions
- The scheduling responsibility of the operating system is primarily limited to processor allocation
- This approach is superior to gang scheduling or dedicated processor assignment for applications that can take advantage of it