

Agenda

✓ Multiprocessor Scheduling

- ✓ Granularity
- ✓ Design Issues
- ✓ Process Scheduling
- ✓ Thread Scheduling

■ Real-Time Scheduling

- Background
- Characteristics of Real-Time Operating Systems
- Real-Time Scheduling
- Deadline Scheduling
- Rate Monotonic Scheduling
- Priority Inversion

Real-Time Systems

- The operating system, and in particular the scheduler, is perhaps the most important component

Examples:

- control of laboratory experiments
- process control in industrial plants
- robotics
- air traffic control
- telecommunications
- military command and control systems

Real-Time Systems

- Correctness of the system depends not only on the logical result of the computation **but also on the time at** which the results are produced.
- Tasks or processes attempt to **control or react to events** that take place in the outside world.
- These events occur in **“real time”** and tasks must be able to keep up with them.

Hard and Soft Real-Time Tasks

Hard real-time task

- one that must meet its deadline.
- otherwise it will cause unacceptable damage or a fatal error to the system.

Soft real-time task

- Has an associated deadline that is desirable but not mandatory
- It still makes sense to schedule and complete the task even if it has passed its deadline.



Periodic and Aperiodic Tasks

■ Periodic tasks

- requirement may be stated as:
 - once per period T
 - exactly T units apart

■ Aperiodic tasks

- has a deadline by which it must finish or start
- may have a constraint on both start and finish time

Characteristics of Real Time Systems

Real-time operating systems have requirements in five general areas:

Determinism

Responsiveness

User control

Reliability

Fail-soft operation

Determinism

- Concerned with how long an operating system delays before acknowledging an interrupt.
- Operations are performed at fixed, **predetermined** times or within predetermined time **intervals**

Determinism

The extent to which an operating system can deterministically satisfy requests depends on:

- the speed with which it can respond to interrupts.
- whether the system has sufficient capacity to handle all requests within the required time.

Responsiveness


- Together with determinism make up the response time to external events
 - critical for real-time systems that must meet timing requirements imposed by individuals, devices, and data flows external to the system.
- Concerned with how long, after acknowledgment, it takes an operating system to service the interrupt.

Responsiveness

Responsiveness includes:


- amount of time required to initially handle the interrupt and begin execution of the interrupt service routine (ISR)
- amount of time required to perform the ISR
- effect of interrupt nesting

User Control



- Generally much broader in a real-time operating system than in ordinary operating systems
- It is essential to allow the user fine-grained control over task priority
- User should be able to distinguish between hard and soft tasks and to specify relative priorities within each class

User Control



May allow user to specify such characteristics as:

- paging or process swapping
- what processes must always be resident in main memory
- what disk transfer algorithms are to be used
- what rights the processes in various priority bands have

Reliability

- More important for real-time systems than non-real time systems
- Real-time systems respond to and control events in real time so loss or degradation of performance may have catastrophic consequences such as:
 - financial loss
 - major equipment damage
 - loss of life



Fail-Soft Operation



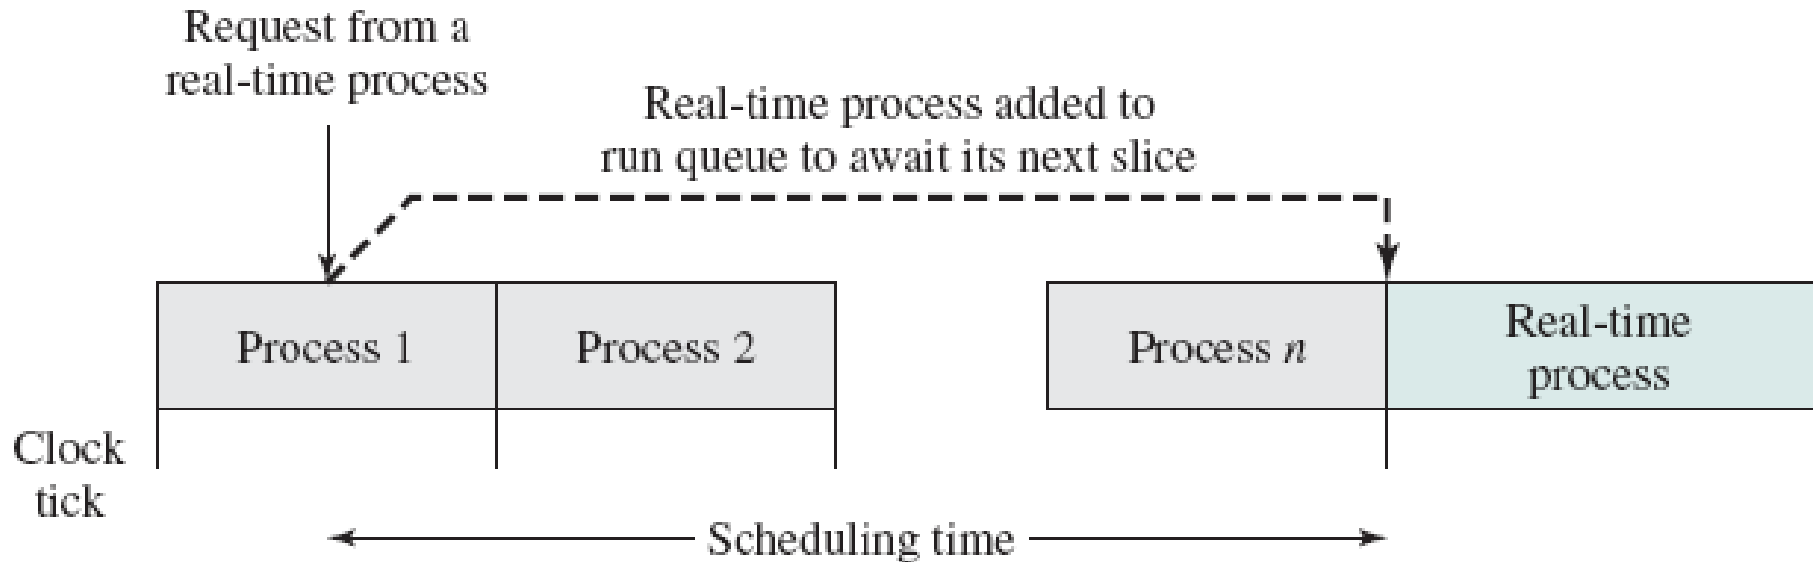
- A characteristic that refers to the ability of a system to fail in such a way as to preserve as much capability and data as possible
- Important aspect is stability
 - a real-time system is stable if the system will meet the deadlines of its most critical, highest-priority tasks even if some less critical task deadlines are not always met

The heart of a real-time system is the short-term task scheduler.

In designing such a scheduler, fairness and minimizing average response time are not important.

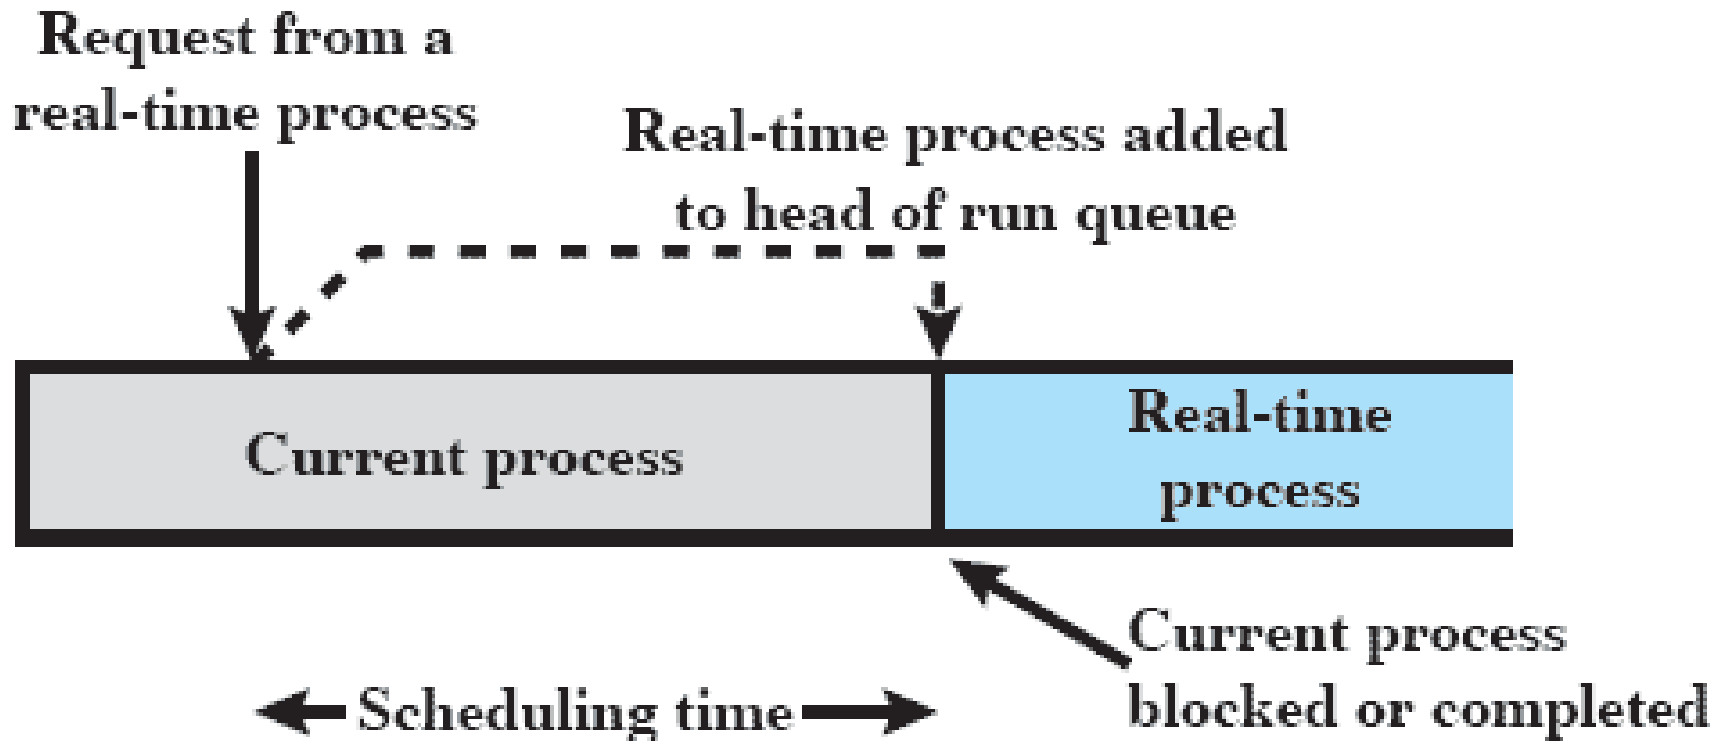
What is important: is that all hard real-time tasks complete (or start) by their deadline and that as many as possible soft real-time tasks also complete (or start) by their deadline.

Scheduling of Real-Time Process



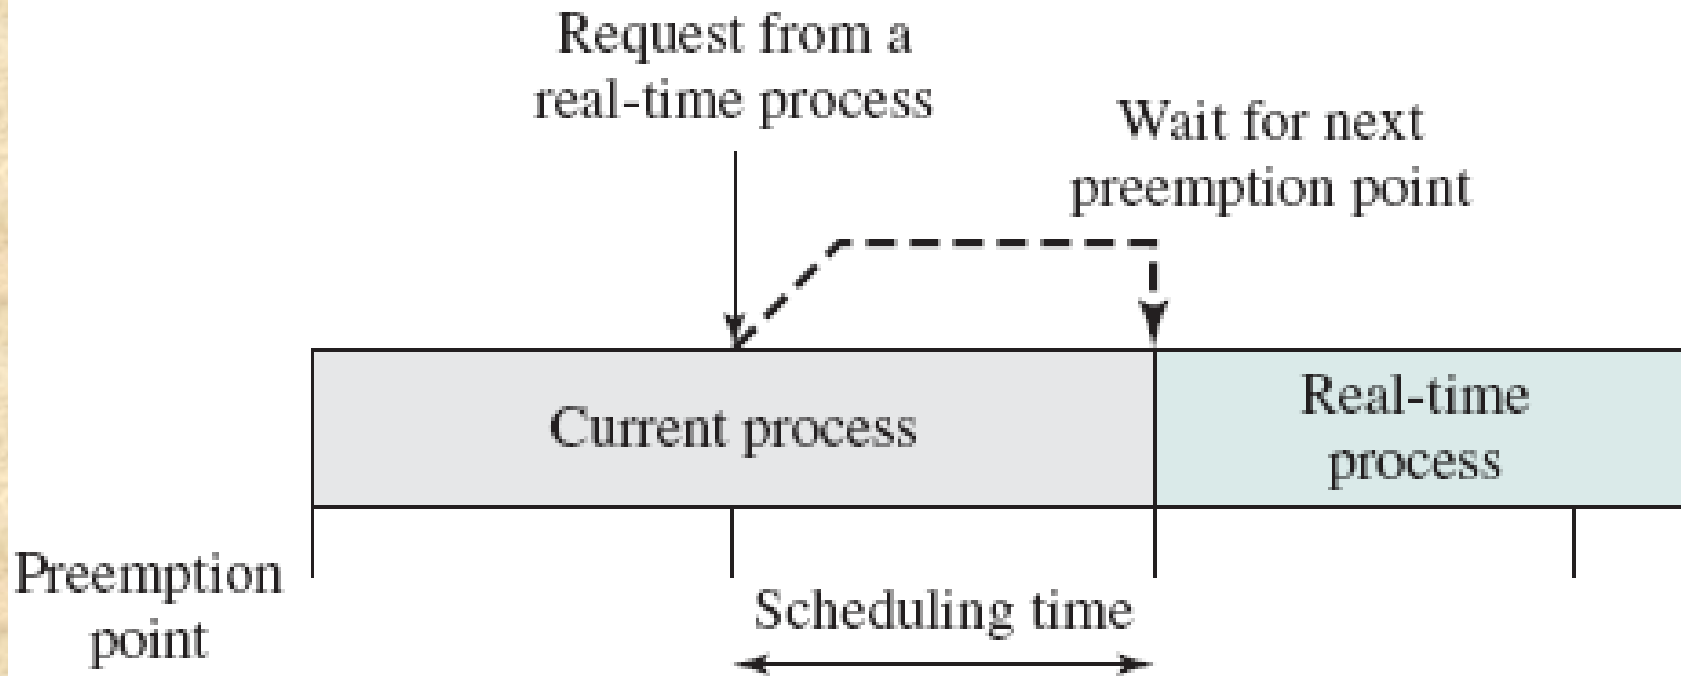
(a) Round-robin preemptive scheduler

In this case, the scheduling time will generally be unacceptable for real-time applications.



(b) Priority-Driven Nonpreemptive Scheduler

This could lead to a delay of several seconds if a slow, low-priority task were executing at a critical time. Again, this approach is not acceptable.



(c) Priority-driven preemptive scheduler on preemption points

A more promising approach is to combine priorities with clock-based interrupts. Preemption points occur at regular intervals. When a preemption point occurs, the currently running task is preempted if a higher-priority task is waiting.

This would include the preemption of tasks that are part of the operating system kernel.

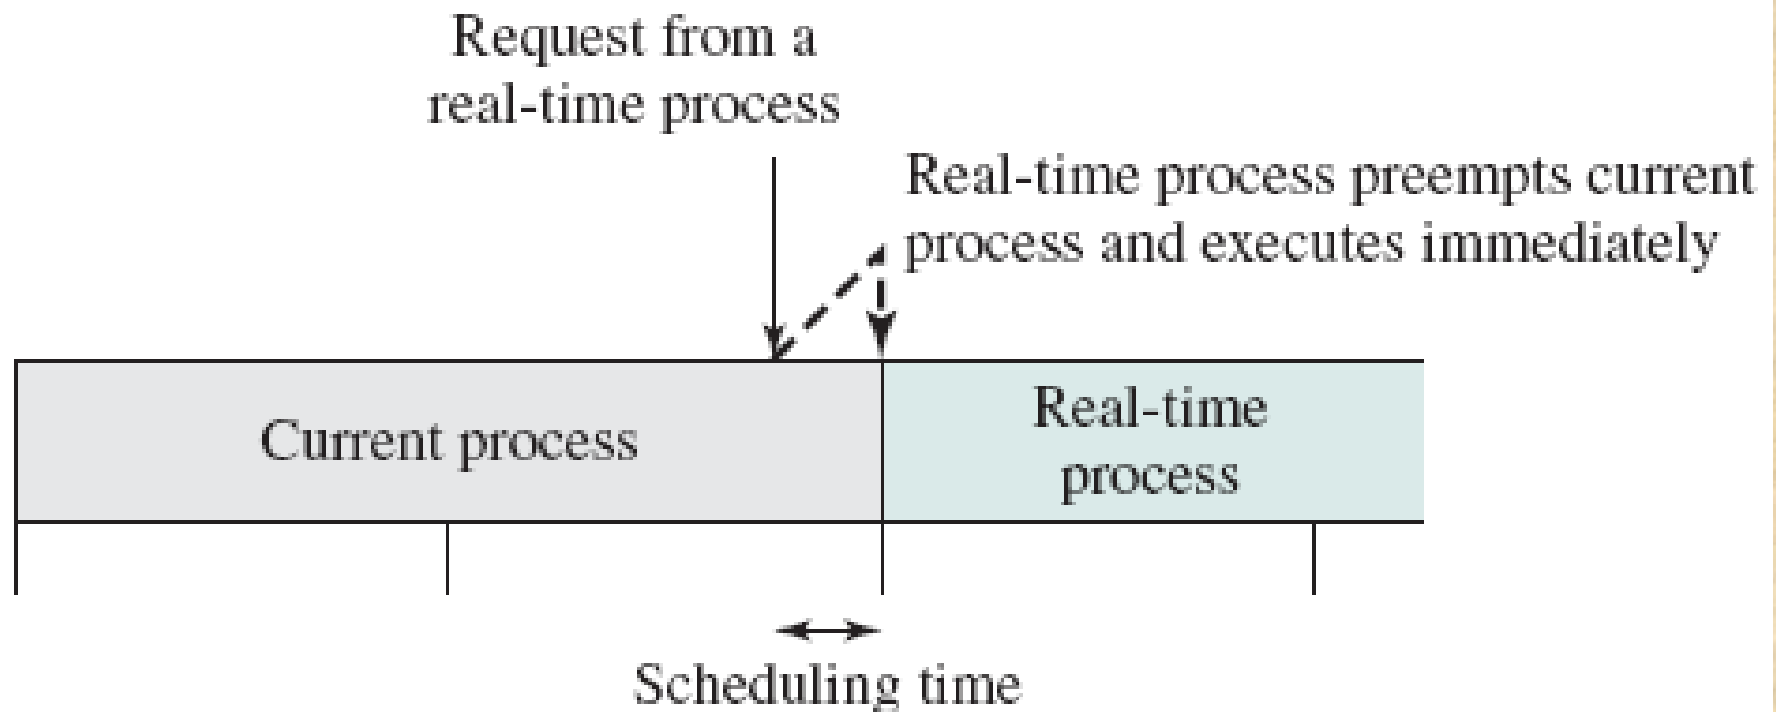
Such a delay may be on the order of several milliseconds (Figure 10.4c).

While this last approach may be adequate for some real-time applications, it will not be enough for more demanding applications.

In those cases, the approach that has been taken is sometimes referred to as immediate preemption.

In this case, the operating system responds to an interrupt almost immediately, unless the system is in a critical-code lockout section.

Scheduling delays for a real-time task can then be reduced to 100 μ s or less.




(d) Immediate preemptive scheduler

Real-Time Scheduling

Scheduling approaches depend on:

- whether a system performs schedulability analysis (if it does, whether it is done statically or dynamically)
- whether the result of the analysis itself produces a scheduler plan according to which tasks are dispatched at run time



Classes of Real-Time Scheduling Algorithms

Static table-driven approaches

- performs a static analysis of feasible schedules of dispatching
- result is a schedule that determines, at run time, when a task must begin execution

Static priority-driven preemptive approaches

- a static analysis is performed but no schedule is drawn up
- analysis is used to assign priorities to tasks so that a traditional priority-driven preemptive scheduler can be used

Classes of Real-Time Scheduling Algorithms

Dynamic planning-based approaches

- feasibility is determined at run time rather than offline prior to the start of execution
- one result of the analysis is a schedule or plan that is used to decide when to dispatch this task

Dynamic best effort approaches

- no feasibility analysis is performed
- system tries to meet all deadlines and aborts any started process whose deadline is missed

Deadline Scheduling



- Real-time operating systems are designed with the objective of starting real-time tasks as rapidly as possible and emphasize rapid interrupt handling and task dispatching
- Real-time applications are generally not concerned with sheer speed but rather with completing (or starting) tasks at the most valuable times
- Priorities provide a basic tool and do not capture the requirement of completion (or initiation) at the most valuable time